

ML with SKLearn

Author: Zuhayr Ali

November 7, 2022

1. Read the Auto data

In [1]:

```
import pandas as pd
df = pd.read_csv('Auto.csv')

print(df.head())
print('\n')
print("dimensions:", df.shape)
```

| | mpg | cylinders | displacement | horsepower | weight | acceleration | year | \ |
|---|------|-----------|--------------|------------|--------|--------------|------|---|
| 0 | 18.0 | 8 | 307.0 | 130 | 3504 | 12.0 | 70.0 | |
| 1 | 15.0 | 8 | 350.0 | 165 | 3693 | 11.5 | 70.0 | |
| 2 | 18.0 | 8 | 318.0 | 150 | 3436 | 11.0 | 70.0 | |
| 3 | 16.0 | 8 | 304.0 | 150 | 3433 | 12.0 | 70.0 | |
| 4 | 17.0 | 8 | 302.0 | 140 | 3449 | NaN | 70.0 | |

| | origin | name |
|---|--------|---------------------------|
| 0 | 1 | chevrolet chevelle malibu |
| 1 | 1 | buick skylark 320 |
| 2 | 1 | plymouth satellite |
| 3 | 1 | amc rebel sst |
| 4 | 1 | ford torino |

dimensions: (392, 9)

2. Data exploration with code

In [2]:

```
print(df.mpg.describe())      # range = (9, 46.6), average = 23.445918
print('\n')
print(df.weight.describe())   # range = (1613, 5140), average = 2977.584184
print('\n')
print(df.year.describe())     # range = (70, 82), average = 76.010256
```

```
count    392.000000
mean      23.445918
std        7.805007
min        9.000000
25%       17.000000
50%       22.750000
75%       29.000000
max       46.600000
Name: mpg, dtype: float64
```

```
count    392.000000
mean    2977.584184
std     849.402560
min    1613.000000
```

```
25%      2225.250000
50%      2803.500000
75%      3614.750000
max       5140.000000
Name: weight, dtype: float64
```

```
count      390.000000
mean        76.010256
std         3.668093
min         70.000000
25%         73.000000
50%         76.000000
75%         79.000000
max         82.000000
Name: year, dtype: float64
```

3. Explore data types

In [3]:

```
print(df.dtypes)
print('\n')

df.cylinders = pd.Categorical(df.cylinders.astype("category").cat.codes)
df['origin'] = pd.Categorical(df.origin)

print(df.dtypes)
```

```
mpg          float64
cylinders     int64
displacement  float64
horsepower    int64
weight        int64
acceleration  float64
year          float64
origin        int64
name          object
dtype: object
```

```
mpg          float64
cylinders     category
displacement  float64
horsepower    int64
weight        int64
acceleration  float64
year          float64
origin        category
name          object
dtype: object
```

4. Deal with NAs

In [4]:

```
df = df.dropna()
print("dimensions:", df.shape)
```

```
dimensions: (389, 9)
```

5. Modify columns

```
In [5]: import numpy as np

df['mpg_high'] = np.where(df.mpg > df.mpg.mean(), 1, 0)
df = df.drop(['mpg', 'name'], axis=1)
print(df.head())
```


| | cylinders | displacement | horsepower | weight | acceleration | year | origin | \ |
|---|-----------|--------------|------------|--------|--------------|------|--------|---|
| 0 | 4 | 307.0 | 130 | 3504 | 12.0 | 70.0 | 1 | |
| 1 | 4 | 350.0 | 165 | 3693 | 11.5 | 70.0 | 1 | |
| 2 | 4 | 318.0 | 150 | 3436 | 11.0 | 70.0 | 1 | |
| 3 | 4 | 304.0 | 150 | 3433 | 12.0 | 70.0 | 1 | |
| 6 | 4 | 454.0 | 220 | 4354 | 9.0 | 70.0 | 1 | |

| | mpg_high |
|---|----------|
| 0 | 0 |
| 1 | 0 |
| 2 | 0 |
| 3 | 0 |
| 6 | 0 |


6. Data exploration with graphs

```
In [6]: import seaborn as sns


mpg_cat = sns.catplot(x='mpg_high', data=df) # I learned nothing I didn't already know
```



```
In [7]: df_rel = sns.relplot(x='horsepower', y='weight', data=df, hue='mpg_high') # Smaller and
```



In [8]: `df_box = sns.boxplot(x='mpg_high', y='weight', data=df) # Further reinforces that small`



7. Train/test split

In [9]: `from sklearn.model_selection import train_test_split`

```
X = df.loc[:, ['cylinders', 'displacement', 'horsepower', 'weight', 'acceleration', 'year']]
y = df.mpg_high

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=1)

print('train size:', X_train.shape)
print('test size:', X_test.shape)
```

train size: (311, 7)
test size: (78, 7)

8. Logistic Regression

```
In [10]: from sklearn.linear_model import LogisticRegression

log_clf = LogisticRegression(solver='lbfgs')
log_clf.fit(X_train, y_train)
print('score:', log_clf.score(X_train, y_train))

log_pred = log_clf.predict(X_test)

from sklearn.metrics import classification_report
print(classification_report(y_test, log_pred))
```

```
score: 0.9067524115755627
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.98 | 0.80 | 0.88 | 50 |
| 1 | 0.73 | 0.96 | 0.83 | 28 |
| accuracy | | | 0.86 | 78 |
| macro avg | 0.85 | 0.88 | 0.85 | 78 |
| weighted avg | 0.89 | 0.86 | 0.86 | 78 |

9. Decision Tree

```
In [11]: from sklearn.tree import DecisionTreeClassifier

dt_clf = DecisionTreeClassifier()
dt_clf.fit(X_train, y_train)
print('score:', dt_clf.score(X_train, y_train))

dt_pred = dt_clf.predict(X_test)

print(classification_report(y_test, dt_pred))
```

```
score: 1.0
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.92 | 0.88 | 0.90 | 50 |
| 1 | 0.80 | 0.86 | 0.83 | 28 |
| accuracy | | | 0.87 | 78 |
| macro avg | 0.86 | 0.87 | 0.86 | 78 |
| weighted avg | 0.87 | 0.87 | 0.87 | 78 |

10. Neural Network

```
In [12]: from sklearn import preprocessing

scaler = preprocessing.StandardScaler().fit(X_train)

X_train_scaled = scaler.transform(X_train)
X_test_scaled = scaler.transform(X_test)

from sklearn.neural_network import MLPClassifier
```

```
nn_clf = MLPClassifier(solver='lbfgs', hidden_layer_sizes=(5, 2), max_iter=500, random_
nn_clf.fit(X_train_scaled, y_train)

nn_pred = nn_clf.predict(X_test_scaled)

from sklearn.metrics import accuracy_score
print('accuracy = ', accuracy_score(y_test, nn_pred))
```

```
accuracy = 0.8717948717948718
```

In []: