

Image Classification

Author: Zuhayr Ali

December 4, 2022

Intro

This notebook will explore using deep learning to classify images from the CIFAR100 dataset. The CIFAR100 is a 60,000 RGB dataset of 32x32 images divided into 20 superclasses based on categories such as plants, animals, humans, buildings, objects, and vehicles; each superclass has 5 subclasses of specific entities belonging to each category, bringing a total of 100 classes. Due to confusion and time constraints the CIFAR100 dataset was loaded from within Keras as one of its sample datasets, selecting to classify by the subclasses. The goal is to accurately classify an image into a subclass - ergo it must predict both the category of subject matter and the specific identity of the subject matter.

Load dataset

```

In [1]: import tensorflow as tf
import numpy as np
import keras

batch_size = 128
num_classes = 100
epochs = 20

(X_train, y_train), (X_test, y_test) = tf.keras.datasets.cifar100.load_data(label

y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

X_train, X_test = X_train / 255.0, X_test / 255.0

print(X_train[:5])
print('\n')
print("dimensions:", X_train.shape)

print('\n')
print(X_test[:5])
print('\n')
print("dimensions:", X_test.shape)

```

```

[[[1.         1.         1.         ]
  [1.         1.         1.         ]
  [1.         1.         1.         ]
  ...
  [0.76470588 0.80392157 0.75686275]
  [0.83137255 0.87843137 0.8         ]
  [0.71372549 0.76078431 0.65490196]]

[[1.         1.         1.         ]
  [0.99607843 0.99607843 0.99607843]
  [0.99607843 0.99607843 0.99607843]
  ...
  [0.66666667 0.69019608 0.58823529]
  [0.63137255 0.65882353 0.50980392]
  [0.57254902 0.60392157 0.44313725]]

[[1.         1.         1.         ]
  [0.99607843 0.99607843 0.99607843]
  [1.         1.         1.         ]

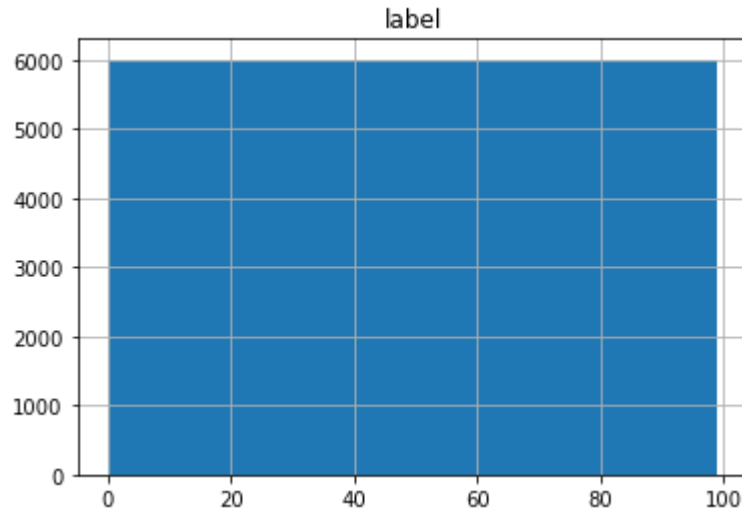
```

Class distribution

```
In [2]: import pandas as pd

df = pd.DataFrame(np.append(np.argmax(y_train, axis=1), np.argmax(y_test, axis=1))
hist = df.hist()
print(hist)
```

```
[[<AxesSubplot:title={'center':'label'}>]]
```



Sequential model

Training

```
In [3]: import keras
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.optimizers import RMSprop

seq_model = Sequential([
    Flatten(input_shape=(32,32,3)),
    Dense(512, activation='relu'),
    Dropout(0.2),
    Dense(512, activation='relu'),
    Dropout(0.2),
    Dense(num_classes, activation='softmax'),
])

seq_model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
flatten (Flatten)	(None, 3072)	0
dense (Dense)	(None, 512)	1573376
dropout (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 512)	262656
dropout_1 (Dropout)	(None, 512)	0
dense_2 (Dense)	(None, 100)	51300
=====		
Total params: 1,887,332		
Trainable params: 1,887,332		
Non-trainable params: 0		

Testing

```
In [4]: seq_model.compile(loss='binary_crossentropy',
                        optimizer=RMSprop(),
                        metrics=['accuracy'])

seq_history = seq_model.fit(X_train, y_train,
                            batch_size=batch_size,
                            epochs=epochs,
                            verbose=1,
                            validation_data=(X_test, y_test))
```

Epoch 1/20

391/391 [=====] - 6s 15ms/step - loss: 0.0673 - accuracy: 0.0417 - val_loss: 0.0560 - val_accuracy: 0.0849

Epoch 2/20

391/391 [=====] - 6s 14ms/step - loss: 0.0527 - accuracy: 0.0913 - val_loss: 0.0498 - val_accuracy: 0.1141

Epoch 3/20

391/391 [=====] - 6s 15ms/step - loss: 0.0493 - accuracy: 0.1127 - val_loss: 0.0483 - val_accuracy: 0.1410

Epoch 4/20

391/391 [=====] - 6s 14ms/step - loss: 0.0483 - accuracy: 0.1282 - val_loss: 0.0474 - val_accuracy: 0.1417

Epoch 5/20

391/391 [=====] - 6s 14ms/step - loss: 0.0475 - accuracy: 0.1393 - val_loss: 0.0468 - val_accuracy: 0.1548

Epoch 6/20

391/391 [=====] - 6s 14ms/step - loss: 0.0471 - accuracy: 0.1471 - val_loss: 0.0457 - val_accuracy: 0.1773

Epoch 7/20

391/391 [=====] - 6s 14ms/step - loss: 0.0467 - accuracy: 0.1549 - val_loss: 0.0453 - val_accuracy: 0.1693

Epoch 8/20

391/391 [=====] - 6s 14ms/step - loss: 0.0465 - accuracy: 0.1602 - val_loss: 0.0455 - val_accuracy: 0.1732

Epoch 9/20

391/391 [=====] - 6s 14ms/step - loss: 0.0464 - accuracy: 0.1628 - val_loss: 0.0459 - val_accuracy: 0.1711

Epoch 10/20

391/391 [=====] - 6s 14ms/step - loss: 0.0462 - accuracy: 0.1676 - val_loss: 0.0451 - val_accuracy: 0.1796

Epoch 11/20

391/391 [=====] - 6s 14ms/step - loss: 0.0461 - accuracy: 0.1667 - val_loss: 0.0452 - val_accuracy: 0.1800

Epoch 12/20

391/391 [=====] - 6s 14ms/step - loss: 0.0460 - accuracy: 0.1702 - val_loss: 0.0460 - val_accuracy: 0.1869

Epoch 13/20

391/391 [=====] - 6s 14ms/step - loss: 0.0459 - accuracy: 0.1744 - val_loss: 0.0458 - val_accuracy: 0.1701

Epoch 14/20

391/391 [=====] - 6s 14ms/step - loss: 0.0459 - accuracy: 0.1745 - val_loss: 0.0448 - val_accuracy: 0.1880

Epoch 15/20

391/391 [=====] - 6s 14ms/step - loss: 0.0459 - accuracy: 0.1766 - val_loss: 0.0453 - val_accuracy: 0.1891

Epoch 16/20

391/391 [=====] - 6s 14ms/step - loss: 0.0458 - accuracy:

```

cy: 0.1796 - val_loss: 0.0447 - val_accuracy: 0.1904
Epoch 17/20
391/391 [=====] - 6s 14ms/step - loss: 0.0457 - accuracy: 0.1780 - val_loss: 0.0446 - val_accuracy: 0.1953
Epoch 18/20
391/391 [=====] - 6s 14ms/step - loss: 0.0457 - accuracy: 0.1783 - val_loss: 0.0442 - val_accuracy: 0.1898
Epoch 19/20
391/391 [=====] - 6s 14ms/step - loss: 0.0458 - accuracy: 0.1773 - val_loss: 0.0446 - val_accuracy: 0.1874
Epoch 20/20
391/391 [=====] - 6s 14ms/step - loss: 0.0457 - accuracy: 0.1802 - val_loss: 0.0451 - val_accuracy: 0.1934

```


Analysis

```

In [5]: import matplotlib.pyplot as plt

plt.plot(seq_history.history['val_accuracy'])
plt.plot(seq_history.history['accuracy'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')
plt.show()

```



```

In [6]: seq_score = seq_model.evaluate(X_test, y_test, verbose=0)
print('Test loss:', seq_score[0])
print('Test accuracy:', seq_score[1])

```

```

Test loss: 0.04513979330658913
Test accuracy: 0.19339999556541443

```

CNN

Training

```
In [7]: from keras import Input
from keras.layers import Conv2D, MaxPooling2D

num_filters = 8
filter_size = 3
pool_size = 2

cnn_model = Sequential(
    [
        Input(shape=(32, 32, 3)),
        Conv2D(32, kernel_size=3, activation="relu"),
        MaxPooling2D(pool_size=2),
        Conv2D(64, kernel_size=3, activation="relu"),
        MaxPooling2D(pool_size=2),
        Flatten(),
        Dropout(0.5),
        Dense(num_classes, activation="softmax"),
    ]
)

cnn_model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d (MaxPooling2D)	(None, 15, 15, 32)	0
conv2d_1 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 6, 6, 64)	0
flatten_1 (Flatten)	(None, 2304)	0
dropout_2 (Dropout)	(None, 2304)	0
dense_3 (Dense)	(None, 100)	230500
=====		
Total params: 249,892		
Trainable params: 249,892		
Non-trainable params: 0		
=====		

Testing

In [8]: `from keras.optimizers import Adam`

```
cnn_model.compile(loss='categorical_crossentropy',
                  optimizer=Adam(),
                  metrics=['accuracy'])

cnn_history = cnn_model.fit(X_train, y_train,
                           batch_size=batch_size,
                           epochs=epochs,
                           verbose=1,
                           validation_data=(X_test, y_test))
```

Epoch 1/20

391/391 [=====] - 13s 33ms/step - loss: 3.9020 - accuracy: 0.1180 - val_loss: 3.4041 - val_accuracy: 0.2091

Epoch 2/20

391/391 [=====] - 13s 32ms/step - loss: 3.2979 - accuracy: 0.2241 - val_loss: 3.0982 - val_accuracy: 0.2689

Epoch 3/20

391/391 [=====] - 13s 33ms/step - loss: 3.0500 - accuracy: 0.2701 - val_loss: 2.9242 - val_accuracy: 0.2996

Epoch 4/20

391/391 [=====] - 13s 32ms/step - loss: 2.9029 - accuracy: 0.2978 - val_loss: 2.8290 - val_accuracy: 0.3249

Epoch 5/20

391/391 [=====] - 12s 32ms/step - loss: 2.7724 - accuracy: 0.3219 - val_loss: 2.7461 - val_accuracy: 0.3330

Epoch 6/20

391/391 [=====] - 13s 32ms/step - loss: 2.6825 - accuracy: 0.3394 - val_loss: 2.6743 - val_accuracy: 0.3489

Epoch 7/20

391/391 [=====] - 13s 32ms/step - loss: 2.6016 - accuracy: 0.3564 - val_loss: 2.6322 - val_accuracy: 0.3518

Epoch 8/20

391/391 [=====] - 13s 32ms/step - loss: 2.5314 - accuracy: 0.3724 - val_loss: 2.5574 - val_accuracy: 0.3698

Epoch 9/20

391/391 [=====] - 13s 33ms/step - loss: 2.4724 - accuracy: 0.3818 - val_loss: 2.5034 - val_accuracy: 0.3777

Epoch 10/20

391/391 [=====] - 13s 33ms/step - loss: 2.4122 - accuracy: 0.3959 - val_loss: 2.4755 - val_accuracy: 0.3898

Epoch 11/20

391/391 [=====] - 13s 33ms/step - loss: 2.3690 - accuracy: 0.4028 - val_loss: 2.4595 - val_accuracy: 0.3877

Epoch 12/20

391/391 [=====] - 13s 33ms/step - loss: 2.3212 - accuracy: 0.4126 - val_loss: 2.4131 - val_accuracy: 0.4009

Epoch 13/20

391/391 [=====] - 13s 33ms/step - loss: 2.2851 - accuracy: 0.4206 - val_loss: 2.4093 - val_accuracy: 0.4016

Epoch 14/20

391/391 [=====] - 13s 33ms/step - loss: 2.2525 - accuracy: 0.4259 - val_loss: 2.3722 - val_accuracy: 0.4102

Epoch 15/20

391/391 [=====] - 13s 33ms/step - loss: 2.2219 - acc


```

uracy: 0.4315 - val_loss: 2.3631 - val_accuracy: 0.4061
Epoch 16/20
391/391 [=====] - 13s 33ms/step - loss: 2.1874 - acc
uracy: 0.4404 - val_loss: 2.3414 - val_accuracy: 0.4177
Epoch 17/20
391/391 [=====] - 13s 33ms/step - loss: 2.1570 - acc
uracy: 0.4452 - val_loss: 2.3322 - val_accuracy: 0.4185
Epoch 18/20
391/391 [=====] - 13s 33ms/step - loss: 2.1395 - acc
uracy: 0.4483 - val_loss: 2.3257 - val_accuracy: 0.4210
Epoch 19/20
391/391 [=====] - 13s 33ms/step - loss: 2.1143 - acc
uracy: 0.4526 - val_loss: 2.3276 - val_accuracy: 0.4238
Epoch 20/20
391/391 [=====] - 13s 34ms/step - loss: 2.0996 - acc
uracy: 0.4576 - val_loss: 2.3260 - val_accuracy: 0.4178


```

Analysis

```

In [9]: plt.plot(cnn_history.history['val_accuracy'])
plt.plot(cnn_history.history['accuracy'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')
plt.show()

```



```

In [10]: cnn_score = cnn_model.evaluate(X_test, y_test, verbose=0)
print('Test loss:', cnn_score[0])
print('Test accuracy:', cnn_score[1])

```

```

Test loss: 2.32599139213562
Test accuracy: 0.41780000925064087

```

Pre-trained model (VGG16)

Training

```
In [11]: from keras import Model
from keras.applications import VGG16
from keras.layers import GlobalAveragePooling2D

X_train, X_test = X_train * 255.0, X_test * 255.0

base_model = VGG16(
    weights='imagenet',
    input_shape=(32, 32, 3),
    include_top=False)

base_model.trainable = False

inputs = Input(shape=(32, 32, 3))
x = base_model(inputs, training=False)
x = GlobalAveragePooling2D()(x)
outputs = Dense(100)(x)
pret_model = Model(inputs, outputs)

pret_model.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
=====		
input_3 (InputLayer)	[(None, 32, 32, 3)]	0
vgg16 (Functional)	(None, 1, 1, 512)	14714688
global_average_pooling2d (GlobalAveragePooling2D)	(None, 512)	0
dense_4 (Dense)	(None, 100)	51300
=====		
Total params: 14,765,988		
Trainable params: 51,300		
Non-trainable params: 14,714,688		
=====		

Testing

```
In [12]: pret_model.compile(loss='categorical_crossentropy',
                             optimizer=Adam(),
                             metrics=['accuracy'])

pret_history = pret_model.fit(X_train, y_train,
                              batch_size=batch_size,
                              epochs=epochs,
                              verbose=1,
                              validation_data=(X_test, y_test))
```

Epoch 1/20

391/391 [=====] - 57s 146ms/step - loss: 8.9790 - accuracy: 0.0255 - val_loss: 8.8312 - val_accuracy: 0.0355

Epoch 2/20

391/391 [=====] - 55s 141ms/step - loss: 8.7044 - accuracy: 0.0389 - val_loss: 8.6733 - val_accuracy: 0.0379

Epoch 3/20

391/391 [=====] - 55s 142ms/step - loss: 8.5454 - accuracy: 0.0428 - val_loss: 8.6058 - val_accuracy: 0.0409

Epoch 4/20

391/391 [=====] - 55s 141ms/step - loss: 8.4708 - accuracy: 0.0444 - val_loss: 8.2620 - val_accuracy: 0.0456

Epoch 5/20

391/391 [=====] - 56s 144ms/step - loss: 8.4461 - accuracy: 0.0499 - val_loss: 8.4399 - val_accuracy: 0.0507

Epoch 6/20

391/391 [=====] - 57s 146ms/step - loss: 8.3813 - accuracy: 0.0562 - val_loss: 8.5920 - val_accuracy: 0.0559

Epoch 7/20

391/391 [=====] - 57s 145ms/step - loss: 8.3800 - accuracy: 0.0604 - val_loss: 8.5528 - val_accuracy: 0.0579

Epoch 8/20

391/391 [=====] - 57s 147ms/step - loss: 8.3893 - accuracy: 0.0650 - val_loss: 8.3382 - val_accuracy: 0.0639

Epoch 9/20

391/391 [=====] - 58s 147ms/step - loss: 8.4302 - accuracy: 0.0680 - val_loss: 8.7121 - val_accuracy: 0.0663

Epoch 10/20

391/391 [=====] - 56s 143ms/step - loss: 8.3344 - accuracy: 0.0683 - val_loss: 8.2156 - val_accuracy: 0.0681

Epoch 11/20

391/391 [=====] - 56s 145ms/step - loss: 8.3469 - accuracy: 0.0703 - val_loss: 8.1798 - val_accuracy: 0.0696

Epoch 12/20

391/391 [=====] - 57s 145ms/step - loss: 8.3264 - accuracy: 0.0735 - val_loss: 8.3917 - val_accuracy: 0.0749

Epoch 13/20

391/391 [=====] - 56s 144ms/step - loss: 8.4681 - accuracy: 0.0772 - val_loss: 8.4942 - val_accuracy: 0.0782

Epoch 14/20

391/391 [=====] - 56s 144ms/step - loss: 8.2913 - accuracy: 0.0793 - val_loss: 8.4078 - val_accuracy: 0.0812

Epoch 15/20

391/391 [=====] - 56s 144ms/step - loss: 8.2810 - accuracy: 0.0826 - val_loss: 8.3883 - val_accuracy: 0.0832

Epoch 16/20

391/391 [=====] - 56s 144ms/step - loss: 8.2514 - accuracy: 0.0850 - val_loss: 8.3883 - val_accuracy: 0.0832

```

racy: 0.0848 - val_loss: 8.4282 - val_accuracy: 0.0839
Epoch 17/20
391/391 [=====] - 56s 144ms/step - loss: 8.3596 - accu
racy: 0.0878 - val_loss: 8.5135 - val_accuracy: 0.0861
Epoch 18/20
391/391 [=====] - 57s 145ms/step - loss: 8.2810 - accu
racy: 0.0889 - val_loss: 8.3794 - val_accuracy: 0.0870
Epoch 19/20
391/391 [=====] - 56s 144ms/step - loss: 8.2581 - accu
racy: 0.0898 - val_loss: 8.4203 - val_accuracy: 0.0864
Epoch 20/20
391/391 [=====] - 56s 143ms/step - loss: 8.4792 - accu
racy: 0.0925 - val_loss: 8.5234 - val_accuracy: 0.0907


```

Analysis

```

In [13]: plt.plot(pret_history.history['val_accuracy'])
plt.plot(pret_history.history['accuracy'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')
plt.show()

```



```

In [14]: pret_score = pret_model.evaluate(X_test, y_test, verbose=0)
print('Test loss:', pret_score[0])
print('Test accuracy:', pret_score[1])

```

```

Test loss: 8.523427963256836
Test accuracy: 0.09070000052452087

```

Discussion

Although none are particularly accurate, the results between models are all significantly different and so might more strongly highlight the different capabilities of each model than had the models been trained on a dataset that's easier to classify. The sequential model's accuracy was in the

middle of the 3 models at 19.34% accuracy. However, despite having the least test loss at less than 0.05, the sequential model had the quickest running epochs at 6 seconds each. The CNN model produced the most accurate predictions at 41.78% accuracy. The test loss and epoch eras were both in the middle of the 3 model at 2.326 and 13 seconds respectively. Although the ratio of accuracy to combined epoch time is the same between the two former models and the CNN model took twice as long, the roughly doubled accuracy of the CNN model makes it much more appealing to me for image classification. I was not able to learn how to preprocess the dataset for the pre-trained VGG16 model, which might explain why the results of using VGG16 were generally the worst. The test loss was the highest at 8.523, but the accuracy was the lowest at 9.07% and the epoch runtime was the longest at between 55 and 58 seconds per epoch. In addition, the former two models reached the maximum accuracy in fewer epochs than the VGG16 model. Perhaps the pre-trained model would have performed far better had I time to properly learn how to implement the model.